

STAT 380: Computational Statistics
2016 FINAL EXAMINATION

Student Solution Document

Time allowed: 3 Hours

This examination paper comprises 16 pages.
Answer ALL questions.
The marks total 70.

Name: Ash Midgley

- Two people decide to play a dice rolling game. Each player has a die that they roll n times. The winner is the player with the highest mean value after n rolls. The two players decide to make it a race, checking the mean value after each roll to determine who is in the lead.

- [2 marks] Write a function that simulates $n = 30$ die rolls for each player. The only input should be the value n . The output should be a vector of n rolls for each player. Test your function for $n = 30$.

```
rolls = function(n) {
  p1 = rep(NA, n)
  p2 = rep(NA, n)
  for (i in 1:n) {
    r1 = sample(1:6, 1)
    p1[i] = r1
    r2 = sample(1:6, 1)
    p2[i] = r2
  }
  return(list(p1, p2))
}
rolls(30)
## [[1]]
## [1] 3 3 3 2 4 3 6 4 4 3 5 6 3 4 4 6 2 6 1 3 3 2 2 1 4 6 6 4 5 6
##
## [[2]]
## [1] 1 1 3 1 3 2 3 2 2 2 3 1 4 1 3 2 5 6 6 2 5 4 2 2 5 2 1 4 6 1
```

- [2 marks] Write a function to find the mean values. The input is a single vector that contains the realization of a sequence of n die rolls. The output is a vector of length n that gives the mean value after each roll (i.e. the i th element of the vector is the mean value of all rolls up to and including the i th roll). Test your function using the input vector x below.

```
x = c(1, 3, 5, 5, 2, 1, 1, 3, 4, 1, 6, 6, 4, 6, 5,
      2, 5, 3, 3, 4, 6, 3, 1, 1, 1, 4, 6, 5, 2, 4)
```

```
findMeans = function(vec) {
  n = length(vec)
  meanvec = rep(NA, n)
  for (i in 1:n) {
    meanvec[i] = mean(vec[1:i])
  }
  return(meanvec)
}
findMeans(x)
## [1] 1.000000 2.000000 3.000000 3.500000 3.200000 2.833333 2.571429 2.
## [9] 2.777778 2.600000 2.909091 3.166667 3.230769 3.428571 3.533333 3.
## [17] 3.529412 3.500000 3.473684 3.500000 3.619048 3.590909 3.478261 3.
## [25] 3.280000 3.307692 3.407407 3.464286 3.413793 3.433333
```

- [2 marks] Write a function that finds who is leading after each roll. The inputs are a vector for each player giving the mean value after every roll. The output is a vector of length n that gives the leader after each roll (you should use the value 1 if player

1 is leading, 2 if player 2 is leading and 3 if it is a tie). Test your function with the input vectors $y1$ and $y2$ below:

```
y1 = c(5, 3.5, 3.67, 3.5, 3.2, 3.33, 3.43, 3.25, 3.56,
       3.5, 3.55, 3.5, 3.46, 3.43, 3.53, 3.69, 3.65, 3.72,
       3.74, 3.85, 3.76, 3.77, 3.7, 3.67, 3.64, 3.69,
       3.63, 3.71, 3.69, 3.73)
y2 = c(6, 3.5, 4.33, 3.5, 4, 3.67, 4, 3.62, 3.33, 3.5,
       3.36, 3.25, 3.46, 3.36, 3.33, 3.25, 3.12, 3, 3.16,
       3.25, 3.29, 3.32, 3.43, 3.33, 3.24, 3.23, 3.26,
       3.29, 3.21, 3.23)
```

```
leader = function(p1, p2) {
  n = length(p1)
  leadervec = rep(NA, n)
  for (i in 1:n) {
    p1M = p1[i]
    p2M = p2[i]
    if (p1M > p2M) {
      leadervec[i] = 1
    } else {
      if (p1M < p2M) {
        leadervec[i] = 2
      } else {
        leadervec[i] = 3
      }
    }
  }
  return(leadervec)
}
leader(y1, y2)
## [1] 2 3 2 3 2 2 2 2 1 3 1 1 3 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

- (d) [2 marks] Write a function that finds the number of lead changes. The function takes as input a vector giving the leader after each roll. The output is a single value that records the number of lead changes throughout the game (this includes changes into and out of a tie). Test your function using the input vector z below:

```
z = c(1, 3, 2, 2, 2, 2, 2, 2, 2, 3, 3, 1, 1, 1, 1,
      1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2)
```

```
leadChanges = function(lvec) {
  n = length(lvec)
  leadC = 0
  prevL = lvec[1]
  for (i in 2:n) {
    currL = lvec[i]
    if (currL != prevL) {
      leadC = leadC + 1
    }
    prevL = currL
  }
}
```

```

    return(leadC)
}
leadChanges(z)
## [1] 5

```

- (e) [3 marks] Write a function that simulates the entire game. The input is the number of rolls n , defaulted to 30. The output should consist of
- A vector for each player that gives the mean value after each roll;
 - A vector giving the leader after each roll;
 - The number of lead changes in the game.

Use your function to simulate a game when $n = 100$.

```

simGame = function(n = 30) {
  rollvals = rolls(n)
  p1rolls = rollvals[[1]]
  p2rolls = rollvals[[2]]
  p1M = findMeans(p1rolls)
  p2M = findMeans(p2rolls)
  leaderVec = leader(p1M, p2M)
  leadChanges = leadChanges(leaderVec)
  return(list(p1M, p2M, leaderVec, leadChanges))
}
out = simGame(100)
out
## [[1]]
## [1] 5.000000 4.000000 4.666667 3.750000 3.400000 3.333333 3.000000 3
## [9] 3.333333 3.100000 3.181818 3.333333 3.461538 3.285714 3.200000 3
## [17] 3.235294 3.111111 3.210526 3.300000 3.380952 3.363636 3.434783 3
## [25] 3.440000 3.461538 3.370370 3.392857 3.379310 3.400000 3.387097 3
## [33] 3.363636 3.352941 3.342857 3.361111 3.405405 3.421053 3.410256 3
## [41] 3.463415 3.428571 3.395349 3.431818 3.466667 3.413043 3.382979 3
## [49] 3.306122 3.320000 3.372549 3.384615 3.339623 3.333333 3.363636 3
## [57] 3.315789 3.362069 3.406780 3.416667 3.442623 3.483871 3.444444 3
## [65] 3.430769 3.424242 3.417910 3.441176 3.449275 3.471429 3.450704 3
## [73] 3.479452 3.472973 3.493333 3.460526 3.428571 3.435897 3.417722 3
## [81] 3.395062 3.414634 3.409639 3.416667 3.447059 3.430233 3.459770 3
## [89] 3.460674 3.488889 3.494505 3.510870 3.483871 3.457447 3.452632 3
## [97] 3.453608 3.479592 3.505051 3.520000
##
## [[2]]
## [1] 2.000000 4.000000 4.666667 4.500000 4.400000 4.666667 4.142857 4
## [9] 4.444444 4.300000 4.090909 4.083333 3.923077 3.714286 3.866667 3
## [17] 3.941176 3.833333 3.736842 3.800000 3.904762 3.818182 3.782609 3
## [25] 3.760000 3.769231 3.740741 3.785714 3.689655 3.666667 3.580645 3
## [33] 3.484848 3.500000 3.571429 3.555556 3.594595 3.578947 3.589744 3
## [41] 3.560976 3.619048 3.674419 3.636364 3.666667 3.630435 3.638298 3
## [49] 3.673469 3.620000 3.588235 3.557692 3.509434 3.481481 3.436364 3
## [57] 3.491228 3.534483 3.559322 3.566667 3.524590 3.548387 3.587302 3
## [65] 3.615385 3.621212 3.656716 3.661765 3.666667 3.700000 3.704225 3
## [73] 3.671233 3.648649 3.626667 3.657895 3.688312 3.666667 3.670886 3

```


```
## [81] 3.703704 3.707317 3.686747 3.678571 3.682353 3.686047 3.666667 3
## [89] 3.696629 3.666667 3.648352 3.619565 3.645161 3.659574 3.642105 3
## [97] 3.680412 3.693878 3.717172 3.720000
##
## [[3]]
## [1] 1 3 3 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
## [38] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
## [75] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
##
## [[4]]
## [1] 2
```

- (f) [4 marks] Plot a line graph of mean value (y -axis) against the number of rolls (x -axis) using the output from the simulation you conducted in part (e). The data for the two players should be plotted in different colours. Ensure your plot is appropriately labeled.

Note: if you were unable to successfully implement the function in part (e) you should use the mean vectors $y1$ and $y2$ given in part (c).

```
plot(1:100, out[[1]], type = "o", col = "red", main = "Mean roll values v
      ylab = "mean roll values", xlab = "number of rolls")
lines(out[[2]], type = "o", col = "blue")
```

Mean roll values vs. number of rolls



(g) [2 marks] Give a step-by-step description of how you could use the functions specified above to estimate the average number of lead changes in a $n = 100$ game using simulation.

- I would first define the m value and set the corresponding probability vector to size m
- I would then loop from 1 to m
- For each iteration I would simulate the rolls with $n=100$, then use this to find the mean roll vectors for each player, then find the current leader vector and then use this to find the number of lead changes
- Once I have found the lead changes for that iteration, I would input it into the probability vector at the current index
- Once the m iterations are complete and we have filled the probability vector, I would find the mean value of the prob vector which would give us the average number of lead changes in a $n=100$ game

2. The dataset **state** contains information on 50 US states and is made up of built-in datasets *state.abb*, *state.division* and *state.x77*.

```
state = data.frame(state = state.abb, region = state.division,
  state.x77)
```

- (a) [2 marks] We wish to recode the factor variable *region*. Replace both the *Pacific* and *Mountain* regions by a new region: *West*.

```
library(dplyr)
##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
levels(state$region) = c("New England", "Middle Atlantic",
  "South Atlantic", "East South Central", "West South Central",
  "East North Central", "West North Central", "West",
  "West")
```

- (b) [1 mark] Add a variable to **state** called *Density*. It is given as

$$Density = \frac{1000 \times Population}{Area}.$$

```
state = mutate(state, Density = ((1000 * Population)/Area))
```

- (c) [1 mark] Find the median and interquartile range of the murder rate across all 50 states.

```
median(state$Murder)
## [1] 6.85
quantile(state$Murder)
##      0%      25%      50%      75%     100%
##  1.400  4.350  6.850 10.675 15.100
```

- (d) [2 marks] Find the total population of each region.

```
summarise(group_by(state, region), sum(state$Population))
## Source: local data frame [8 x 2]
##
##           region sum(state$Population)
##           (fctr)                (dbl)
## 1      New England                212321
## 2  Middle Atlantic                212321
## 3    South Atlantic                212321
## 4 East South Central                212321
## 5 West South Central                212321
## 6 East North Central                212321
## 7 West North Central                212321
## 8                West                212321
```

- (e) [2 marks] Obtain a new data frame that only features states with *Population* more than 750 and *Frost* less than 170.

```
d2 = filter(state, Population > 750 | Frost < 150)
d2
```

##	state	region	Population	Income	Illiteracy	Life.Exp	Murder
## 1	AL East	South Central	3615	3624	2.1	69.05	15
## 2	AZ	West	2212	4530	1.8	70.55	
## 3	AR West	South Central	2110	3378	1.9	70.66	10
## 4	CA	West	21198	5114	1.1	71.71	10
## 5	CO	West	2541	4884	0.7	72.06	6
## 6	CT	New England	3100	5348	1.1	72.48	3
## 7	DE	South Atlantic	579	4809	0.9	70.06	6
## 8	FL	South Atlantic	8277	4815	1.3	70.66	10
## 9	GA	South Atlantic	4931	4091	2.0	68.54	13
## 10	HI	West	868	4963	1.9	73.60	6
## 11	ID	West	813	4119	0.6	71.87	5
## 12	IL East	North Central	11197	5107	0.9	70.14	10
## 13	IN East	North Central	5313	4458	0.7	70.88	7
## 14	IA West	North Central	2861	4628	0.5	72.56	2
## 15	KS West	North Central	2280	4669	0.6	72.58	4
## 16	KY East	South Central	3387	3712	1.6	70.10	10
## 17	LA West	South Central	3806	3545	2.8	68.76	13
## 18	ME	New England	1058	3694	0.7	70.39	2
## 19	MD	South Atlantic	4122	5299	0.9	70.22	8
## 20	MA	New England	5814	4755	1.1	71.83	3
## 21	MI East	North Central	9111	4751	0.9	70.63	11
## 22	MN West	North Central	3921	4675	0.6	72.96	2
## 23	MS East	South Central	2341	3098	2.4	68.09	12
## 24	MO West	North Central	4767	4254	0.8	70.69	9
## 25	NE West	North Central	1544	4508	0.6	72.60	2
## 26	NH	New England	812	4281	0.7	71.23	3
## 27	NJ	Middle Atlantic	7333	5237	1.1	70.93	5
## 28	NM	West	1144	3601	2.2	70.32	9
## 29	NY	Middle Atlantic	18076	4903	1.4	70.55	10
## 30	NC	South Atlantic	5441	3875	1.8	69.21	11
## 31	OH East	North Central	10735	4561	0.8	70.82	7
## 32	OK West	South Central	2715	3983	1.1	71.42	6
## 33	OR	West	2284	4660	0.6	72.13	4
## 34	PA	Middle Atlantic	11860	4449	1.0	70.43	6
## 35	RI	New England	931	4558	1.3	71.90	2
## 36	SC	South Atlantic	2816	3635	2.3	67.96	11
## 37	TN East	South Central	4173	3821	1.7	70.11	11
## 38	TX West	South Central	12237	4188	2.2	70.90	12
## 39	UT	West	1203	4022	0.6	72.90	4
## 40	VA	South Atlantic	4981	4701	1.4	70.08	9
## 41	WA	West	3559	4864	0.6	71.72	4
## 42	WV	South Atlantic	1799	3617	1.4	69.48	6
## 43	WI East	North Central	4589	4468	0.7	72.48	3
##	HS.Grad	Frost	Area	Density			
## 1	41.3	20	50708	71.290526			
## 2	58.1	15	113417	19.503249			


```
## 3      39.9      65  51945  40.619886
## 4      62.6      20 156361 135.570890
## 5      63.9     166 103766  24.487790
## 6      56.0     139   4862 637.597696
## 7      54.6     103   1982 292.129162
## 8      52.6      11  54090 153.022740
## 9      40.6      60  58073  84.910371
## 10     61.9       0   6425 135.097276
## 11     59.5     126  82677   9.833448
## 12     52.6     127  55748 200.850255
## 13     52.9     122  36097 147.186747
## 14     59.0     140  55941  51.143169
## 15     59.9     114  81787  27.877291
## 16     38.5      95  39650  85.422446
## 17     42.2      12  44930  84.709548
## 18     54.7     161  30920  34.217335
## 19     52.3     101   9891 416.742493
## 20     58.5     103   7826 742.908255
## 21     52.8     125  56817 160.356935
## 22     57.6     160  79289  49.452005
## 23     41.0      50  47296  49.496786
## 24     48.8     108  68995  69.091963
## 25     59.3     139  76483  20.187493
## 26     57.6     174   9027  89.952365
## 27     52.5     115   7521 975.003324
## 28     55.2     120 121412   9.422462
## 29     52.7      82  47831 377.913905
## 30     38.5      80  48798 111.500471
## 31     53.2     124  40975 261.989018
## 32     51.6      82  68782  39.472536
## 33     60.0      44  96184  23.746153
## 34     50.2     126  44966 263.754837
## 35     46.4     127   1049 887.511916
## 36     37.8      65  30225  93.167907
## 37     41.8      70  41328 100.972706
## 38     47.4      35 262134  46.682231
## 39     67.3     137  82096  14.653576
## 40     47.8      85  39780 125.213675
## 41     63.5      32  66570  53.462521
## 42     41.6     100  24070  74.740341
## 43     54.5     149  54464  84.257491
```

- (f) [3 marks] Regress *Income* on *HS.Grad*. Find a 95% confidence interval for the effect of graduation rate on income. Interpret the confidence interval.

```
m1 = lm(state$Income ~ state$HS.Grad)
confint(m1)

##              2.5 %          97.5 %
## (Intercept) 1000.70545 2861.50393
## state$HS.Grad 29.83854  64.48606
```

The confidence interval is between 29.83854 and 64.48606

3. (a) [3 marks] State the difference between a list and a vector. Give an example where a list is preferred over a vector.
- Lists allow storage of multiple individual objects under one object name
 - Vectors are only allowed to be made up of a single class, whereas lists allow objects of different classes
 - Each element in a list can be a vector, a matrix or another list etc
 - Example below shows how vectors don't store vectors properly but lists can
- EXAMPLE:

```
workingList = list(c(1, 2, 3), c("a", "b"))
errorVec = c(c(1, 2, 3), c("a", "b"))
workingList
## [[1]]
## [1] 1 2 3
##
## [[2]]
## [1] "a" "b"
errorVec
## [1] "1" "2" "3" "a" "b"
```

- (b) [4 marks] We have measured the height of 30 kahikatea trees around Otago. We have estimated the mean tree height and used a large-sample confidence interval based on the central limit theorem. Give step-by-step directions for how we could use simulation to estimate the coverage of the large-sample confidence interval.
- Set number of simulations (m)
 - Find true mean based on theta
 - Make a vector with sample sizes across the range
 - store the length of this vec in variable n
 - set a coverage vec to this size
 - Iterate through from 1 to n
 - For each iteration, find the confidence interval
 - Check whether the true mean is within this confidence interval
 - If so add 1 to index of coverage vec, otherwise 0
 - After we have finished iterating, we can view the coverage in a data frame
4. Researchers are interested in the number of faults in rolls of fabric as a function of the length of the roll. The first 5 rows of the data are in the table below.

- (a) [6 marks] Use Poisson regression via the `glm()` function to fit a model of the form:

$$\log(\lambda_i) = \beta_0 + \beta_1 x_i$$

where λ_i is expected number of faults in a roll of length x_i . Provide a brief summary of your model fit, including a plot of the expected number of faults vs roll length. This plot should include a line showing the fitted values under the model.


Faults	Length
6	551
4	651
17	832
9	375
14	715

Table 1: The number of faults in rolls of fabric as a function of the length of the roll

```

faults = c(6, 4, 17, 9, 14)
length = c(551, 651, 832, 375, 715)
fit = glm(faults ~ length, family = poisson)
fit
##
## Call:  glm(formula = faults ~ length, family = poisson)
##
## Coefficients:
## (Intercept)      length
##    1.030585    0.001965
##
## Degrees of Freedom: 4 Total (i.e. Null);  3 Residual
## Null Deviance:      12.11
## Residual Deviance:  7.892  AIC: 32.04
plot(length, faults)
lines(length, fit$fitted.values, col = "red")

```



- (b) [2 marks] Use your parameter estimates from part (a) to estimate θ , the value of x_i when the expected number of faults is exactly 5.

```
logval = log(5)
theta = (logval - beta[1])/beta[2]
## Error in beta[1]: object of type 'closure' is not subsettable
theta
## Error in eval(expr, envir, enclos): object 'theta' not found
```

- (c) [6 marks] Use the parametric bootstrap to find a 95% confidence interval for θ .

```
Q <- 999
bhatboot = matrix(NA, Q, 2)
thetaboot = rep(NA, Q)

for (i in 1:Q) {
  r = rbinom(length(faults), faults, prob = fit$fit)
  fitboot <- glm(r ~ length, family = poisson)
  bhatboot[i, ] = fitboot$coef
  thetaboot[i] = (logval - bhatboot[i, 1])/bhatboot[i,
```

```

      2]
}
## Warning in rbinom(length(faults), faults, prob = fit$fit):
NAs produced
## Warning in glm.fit(x = structure(numeric(0), .Dim = c(0L,
2L), .Dimnames = list(: no observations informative at iteration
1
## Warning: glm.fit: algorithm did not converge
## Error in fit$qr[1L:nr, 1L:nvars]: subscript out of bounds
ci = 2 * mean(thetaboot) - quantile(thetaboot, c(0.975,
0.025))
## Error in quantile.default(thetaboot, c(0.975, 0.025)):
missing values and NaN's not allowed if 'na.rm' is FALSE
ci
## Error in eval(expr, envir, enclos): object 'ci' not found

```

- (d) [2 marks] Briefly explain the difference between parametric and non-parametric bootstrapping.

Parametric:

- Assume some parametric model for the underlying population F_{θ}
- Estimate the parameters of this model
- Resample from this model using the estimated parameters to estimate uncertainty

Non-parametric:

- Make few assumptions about the underlying distribution
- Use \hat{F} as an estimator for F
- \hat{F} is non-parametric estimator for F

- (e) [5 marks] Use JAGS implemented through R using R2jags to redo the analysis described above. Use normal $\mathcal{N}(0, 10000)$ (mean and variance) prior distributions for the parameters β_0 and β_1 . As output, report traceplots for the parameters of the model and report statistics that provide posterior summaries.

```

library("R2jags")
## Loading required package: rjags
## Loading required package: coda
## Linked to JAGS 4.2.0
## Loaded modules: basemod, bugs
##
## Attaching package: 'R2jags'
## The following object is masked from 'package:coda':
##
## traceplot
model = function() {
  for (i in 1:n) {
    logit(p[i]) <- b[1] + b[2] * length[i]
  }
  for (i in 1:2) {

```

```

        b[i] ~ dnorm(0, 1e-04)
    }
    theta <- (logit(5) - b[i])/b[2]
}
data = c("Faults", "Length", "n")

inits = function() {
    list(b = rnorm(2, mean = beta, sd = 0.1))
}

params = c("b", "theta")
jagsfit = jags(data = data, inits, parameters.to.save = params,
               n.iter = 1e+05, model.file = model)
## module glm loaded
## Error in rnorm(2, mean = beta, sd = 0.1): invalid arguments
traceplot(as.mcmc(jagsfit), ask = F)
## Error in traceplot(as.mcmc(jagsfit), ask = F): error in
evaluating the argument 'x' in selecting a method for function
'traceplot': Error in as.mcmc(jagsfit) : object 'jagsfit'
not found
jfsum = jagsfit$BUGS$sum
## Error in eval(expr, envir, enclos): object 'jagsfit'
not found

```

- (f) [2 marks] From your analysis in part (e), what is a 95% credible interval for θ ?

```

theta95 = jfsum[4, c(3, 7)]
## Error in eval(expr, envir, enclos): object 'jfsum' not
found

```

5. The pdf for a random variable with support on $(0,1)$ is given by

$$f_X(x) = \begin{cases} 12(x^2 - x^3) & 0 \leq x \leq 1 \\ 0 & \text{otherwise} \end{cases}$$

and with cumulative density function (cdf) given by

$$F_X(x) = \begin{cases} 0 & x < 0 \\ 4x^3 - 3x^4 & 0 \leq x \leq 1 \\ 1 & x > 1 \end{cases}$$

This distribution has a mode at $2/3$.

- (a) [4 marks] Write an R function called `dfx` that returns the pdf of the above distribution for any x .

```

dfx = function(x) {
    brk = x <= 1
    result = brk * (12 * (x^2 - x^3))
    return(result)
}


```

- (b) [4 marks] Write an R function called `rfx` that uses rejection sampling to generate exactly n samples from the above distribution with the number of samples defaulted to 1.

```
rfx = function(n = 1) {
  M = 1/sqrt(2 * pi)
  u = runif(n)
  x = runif(n, -1.5, 1.5)
  keep = which(u < dnorm(x)/M)
  y = x[keep]
  return(y)
}
```

- (c) [3 marks] Use your functions to generate 10000 samples from the distribution with pdf $f_X(x)$. Display them using an appropriately labelled histogram with 25 breaks. You should overlay in red the corresponding probability density function.

```
x = rfx(1000)
hist(x, breaks = 25, prob = TRUE)
curve(dfx(x), add = T, col = "red")
```



- (d) [1 mark] Based on your simulations what is the expected value of our random variable $\mathbb{E}[X]$?

0

[1] 0